

## 9 most useful functions for PySpark DataFrame

### Pyspark

PySpark is a data analytics tool created by Apache Spark Community for using Python along with Spark. It allows us to work with RDD (Resilient Distributed Dataset) and DataFrames in Python. PySpark has numerous features that make it such an amazing framework and when it comes to deal with the huge amount of data PySpark provides us fast and Real-time processing, flexibility, in-memory computation, and various other features. It is a Python library to use Spark which combines the simplicity of Python language with the efficiency of Spark.

### Pyspark DataFrame

A DataFrame is a distributed collection of data in rows under named columns. In simple terms, we can say that it is the same as a table in a Relational database or an Excel sheet with Column headers. DataFrames are mainly designed for processing a large-scale collection of structured or semi-structured data.

In this article, we'll discuss functions of PySpark that are most useful and essential to perform efficient data analysis of structured data.

We are using Google Colab as the IDE for this data analysis.

We first need to install PySpark in Google Colab. After that, we will import the `pyspark.sql` module and create a `SparkSession` which will be an entry point of Spark SQL API.

```
#installing pyspark
!pip install pyspark
```

```
#importing pyspark
import pyspark

#importing sparksession
from pyspark.sql import SparkSession

#creating a sparksession object and providing appName
spark=SparkSession.builder.appName("pysparkdf").getOrCreate()
```

This SparkSession object will interact with the functions and methods of Spark SQL. Now, let's create a Spark DataFrame by reading a CSV file. We will be using simple dataset i.e. [Nutrition Data on 80 Cereal products](#) available on Kaggle.

```
#creating a dataframe using spark object by reading csv file
df = spark.read.option("header", "true").csv("/content/cereal.csv")
```

```
#show df created top 10 rows
df.show(10)
```

	name	mfr	type	calories	protein	fat	sodium	fiber	carbo	sugars	potass	vitamin
	100% Bran	N	C	70	4	1	130	10	5	6	280	2
	100% Natural Bran	Q	C	120	3	5	15	2	8	8	135	
	All-Bran	K	C	70	4	1	260	9	7	5	320	2
	All-Bran with Ext...	K	C	50	4	0	140	14	8	0	330	2
	Almond Delight	R	C	110	2	2	200	1	14	8	-1	2
	Apple Cinnamon Ch...	G	C	110	2	2	180	1.5	10.5	10	70	2
	Apple Jacks	K	C	110	2	0	125	1	11	14	30	2
	Basic 4	G	C	130	3	2	210	2	18	8	100	2
	Bran Chex	R	C	90	2	1	200	4	15	6	125	2
	Bran Flakes	P	C	90	3	0	210	5	13	5	190	2

only showing top 10 rows

This is the DataFrame we are using for Data analysis. Now, let's print the schema of the DataFrame to know more about the dataset.

```
df.printSchema()
```

```

root
|-- name: string (nullable = true)
|-- mfr: string (nullable = true)
|-- type: string (nullable = true)
|-- calories: string (nullable = true)
|-- protein: string (nullable = true)
|-- fat: string (nullable = true)
|-- sodium: string (nullable = true)
|-- fiber: string (nullable = true)
|-- carbo: string (nullable = true)
|-- sugars: string (nullable = true)
|-- potass: string (nullable = true)
|-- vitamins: string (nullable = true)
|-- shelf: string (nullable = true)
|-- weight: string (nullable = true)
|-- cups: string (nullable = true)
|-- rating: string (nullable = true)

```

The DataFrame consists of 16 features or columns. Each column contains string-type values.

Let's get started with the functions:

- select():** The select function helps us to display a subset of selected columns from the entire dataframe we just need to pass the desired column names. Let's print any three columns of the dataframe using select().

```
df.select('name', 'mfr', 'rating').show(10)
```

```

+-----+-----+-----+
|          name|mfr|   rating|
+-----+-----+-----+
|    100% Bran| N|68.402973|
| 100% Natural Bran| Q|33.983679|
|    All-Bran| K|59.425505|
|All-Bran with Ext...| K|93.704912|
|   Almond Delight| R|34.384843|
|Apple Cinnamon Ch...| G|29.509541|
|   Apple Jacks| K|33.174094|
|      Basic 4| G|37.038562|
|   Bran Chex| R|49.120253|
|   Bran Flakes| P|53.313813|
+-----+-----+-----+
only showing top 10 rows

```

In the output, we got the subset of the dataframe with three columns name, mfr, rating.

- withColumn():** The withColumn function is used to manipulate a column or to create a new column with the existing column. It is a transformation function, we can also change the datatype of any existing column.

In the DataFrame schema, we saw that all the columns are of string type. Let's change the data type of calorie column to an integer.

```
df.withColumn("Calories",df['calories'].cast("Integer")).printSchema()
```

```
root
|-- name: string (nullable = true)
|-- mfr: string (nullable = true)
|-- type: string (nullable = true)
|-- Calories: integer (nullable = true)
|-- protein: string (nullable = true)
|-- fat: string (nullable = true)
|-- sodium: string (nullable = true)
|-- fiber: string (nullable = true)
|-- carbo: string (nullable = true)
|-- sugars: string (nullable = true)
|-- potass: string (nullable = true)
|-- vitamins: string (nullable = true)
|-- shelf: string (nullable = true)
|-- weight: string (nullable = true)
|-- cups: string (nullable = true)
|-- rating: string (nullable = true)
```

In the schema, we can see that the Datatype of calories column is changed to the integer type.

•**groupBy():** The groupBy function is used to collect the data into groups on DataFrame and allows us to perform aggregate functions on the grouped data. This is a very common data analysis operation similar to groupBy clause in SQL.

Let's find out the count of each cereal present in the dataset.

```
df.groupBy("name", "calories").count().show()
```

name	calories	count
Basic 4	130	1
Cocoa Puffs	110	1
Strawberry Fruit ...	90	1
Great Grains Pecan	120	1
Wheat Chex	100	1
MuesliX Crispy Blend	160	1
Raisin Nut Bran	100	1
Honey Nut Cheerios	110	1
Corn Flakes	100	1
Trix	110	1
Grape Nuts Flakes	100	1
Muesli Raisins; P...	150	1
Fruity Pebbles	110	1
Shredded Wheat 'n...	90	1
Post Nat. Raisin ...	120	1
Total Raisin Bran	140	1
Corn Pops	110	1
Cracklin' Oat Bran	110	1
Cinnamon Toast Cr...	120	1
Kix	110	1

only showing top 20 rows

•**orderBy():** The orderBy function is used to sort the entire dataframe based on the particular column of the dataframe. It sorts the rows of the dataframe according to column values. By default, it sorts in ascending order.

name	mfr	type	calories	protein	fat	sodium	fiber	carbo	sugars	potass	vitamin
100% Bran	N	C	70	4	1	130	10	5	6	280	2
100% Natural Bran	Q	C	120	3	5	15	2	8	8	135	2
All-Bran	K	C	70	4	1	260	9	7	5	320	2
All-Bran with Ext...	K	C	50	4	0	140	14	8	0	330	2
Almond Delight	R	C	110	2	2	200	1	14	8	-1	2
Apple Cinnamon Ch...	G	C	110	2	2	180	1.5	10.5	10	70	2
Apple Jacks	K	C	110	2	0	125	1	11	14	30	2
Basic 4	G	C	130	3	2	210	2	18	8	100	2
Bran Chex	R	C	90	2	1	200	4	15	6	125	2
Bran Flakes	P	C	90	3	0	210	5	13	5	190	2

only showing top 10 rows

This is the DataFrame we are using for Data analysis. Now, let's print the schema of the DataFrame to know more about the dataset.

```
df.printSchema()
```

```

root
|-- name: string (nullable = true)
|-- mfr: string (nullable = true)
|-- type: string (nullable = true)
|-- calories: string (nullable = true)
|-- protein: string (nullable = true)
|-- fat: string (nullable = true)
|-- sodium: string (nullable = true)
|-- fiber: string (nullable = true)
|-- carbo: string (nullable = true)
|-- sugars: string (nullable = true)
|-- potass: string (nullable = true)
|-- vitamins: string (nullable = true)
|-- shelf: string (nullable = true)
|-- weight: string (nullable = true)
|-- cups: string (nullable = true)
|-- rating: string (nullable = true)

```

The DataFrame consists of 16 features or columns. Each column contains string-type values.

•**split():** The split() is used to split a string column of the dataframe into multiple columns. This function is applied to the dataframe with the help of withColumn() and select().

The name column of the dataframe contains values in two string words. Let's split the name column into two columns from space between two strings.

```

from pyspark.sql.functions import split

df1 = df.withColumn('Name1', split(df['name'], " ").getItem(0))
        .withColumn('Name2', split(df['name'], " ").getItem(1))

df1.select("name", "Name1", "Name2").show()

```

name	Name1	Name2
100% Bran	100%	Bran
100% Natural Bran	100%	Natural
All-Bran	All-Bran	null
All-Bran with Ext...	All-Bran	with
Almond Delight	Almond	Delight
Apple Cinnamon Ch...	Apple	Cinnamon
Apple Jacks	Apple	Jacks
Basic 4	Basic	4
Bran Chex	Bran	Chex
Bran Flakes	Bran	Flakes
Cap'n'Crunch	Cap'n'Crunch	null
Cheerios	Cheerios	null
Cinnamon Toast Cr...	Cinnamon	Toast
Clusters	Clusters	null
Cocoa Puffs	Cocoa	Puffs
Corn Chex	Corn	Chex
Corn Flakes	Corn	Flakes
Corn Pops	Corn	Pops
Count Chocula	Count	Chocula
Cracklin' Oat Bran	Cracklin'	Oat

only showing top 20 rows

In this output, we can see that the name column is split into columns.

- **lit():** The lit function is used to add a new column to the dataframe that contains literals or some constant value.

Let's add a column "intake quantity" which contains a constant value for each of the cereals along with the respective cereal name.

```
from pyspark.sql.functions import lit

df2 = df.select(col("name"),lit("75 gm").alias("intake quantity"))
df2.show()
```



name	intake quantity
100% Bran	75 gm
100% Natural Bran	75 gm
All-Bran	75 gm
All-Bran with Ext...	75 gm
Almond Delight	75 gm
Apple Cinnamon Ch...	75 gm
Apple Jacks	75 gm
Basic 4	75 gm
Bran Chex	75 gm
Bran Flakes	75 gm
Cap'n'Crunch	75 gm
Cheerios	75 gm
Cinnamon Toast Cr...	75 gm
Clusters	75 gm
Cocoa Puffs	75 gm
Corn Chex	75 gm
Corn Flakes	75 gm
Corn Pops	75 gm
Count Chocula	75 gm
Cracklin' Oat Bran	75 gm

only showing top 20 rows

In the output, we can see that a new column is created “intak quantity” that contains the in-take a quantity of each cereal.

**when():** The when the function is used to display the output based on the particular condition. It evaluates the condition provided and then returns the values accordingly. It is a SQL function that supports PySpark to check multiple conditions in a sequence and return the value. This function similarly works as if-then-else and switch statements.

Let’s see the cereals that are rich in vitamins.

```
from pyspark.sql.functions import when
```

```
df.select("name", when(df.vitamins >= "25", "rich in vitamins")).show()
```



name	CASE WHEN (vitamins >= 25) THEN rich in vitamins END
100% Bran	rich in vitamins
100% Natural Bran	null
All-Bran	rich in vitamins
All-Bran with Ext...	rich in vitamins
Almond Delight	rich in vitamins
Apple Cinnamon Ch...	rich in vitamins
Apple Jacks	rich in vitamins
Basic 4	rich in vitamins
Bran Chex	rich in vitamins
Bran Flakes	rich in vitamins
Cap'n'Crunch	rich in vitamins
Cheerios	rich in vitamins
Cinnamon Toast Cr...	rich in vitamins
Clusters	rich in vitamins
Cocoa Puffs	rich in vitamins
Corn Chex	rich in vitamins
Corn Flakes	rich in vitamins
Corn Pops	rich in vitamins
Count Chocula	rich in vitamins
Cracklin' Oat Bran	rich in vitamins

only showing top 20 rows

•**filter():** The filter function is used to filter data in rows based on the particular column values. For example, we can filter the cereals which have calories equal to 100.

```
from pyspark.sql.functions import filter
```

```
df.filter(df.calories == "100").show()
```

	name	mfr	type	calories	protein	fat	sodium	fiber	carbo	sugars	potass	vitamins
	Corn Flakes	K	C	100	2	0	290	1	21	2	35	25
	Cream of Wheat (Q...	N	H	100	3	0	80	1	21	0	-1	0
	Crispy Wheat & Ra...	G	C	100	2	1	140	2	11	10	120	25
	Double Chex	R	C	100	2	0	190	1	18	5	80	25
	Frosted Mini-Wheats	K	C	100	3	0	0	3	14	7	100	25
	Golden Crisp	P	C	100	2	0	45	0	11	15	40	25
	Grape Nuts Flakes	P	C	100	3	1	140	3	15	5	85	25
	Life	Q	C	100	4	2	150	2	12	6	95	25
	Maypo	A	H	100	4	1	0	0	16	3	95	25
	Multi-Grain Cheerios	G	C	100	2	1	220	2	15	6	90	25
	Product 19	K	C	100	3	0	320	1	20	3	45	100
	Quaker Oat Squares	Q	C	100	4	1	135	2	14	6	110	25
	Quaker Oatmeal	Q	H	100	5	2	0	2.7	-1	-1	110	0
	Raisin Nut Bran	G	C	100	3	2	140	2.5	10.5	8	140	25
	Total Whole Grain	G	C	100	3	1	200	3	16	3	110	100
	Wheat Chex	R	C	100	3	1	230	3	17	3	115	25
	Wheaties	G	C	100	3	1	200	3	17	3	110	25

In this output, we can see that the data is filtered according to the cereals which have 100 calories.

**isNull()/isNotNull():** These two functions are used to find out if there is any null value present in the DataFrame. It is the most essential function for data processing. It is the major tool used for data cleaning.

Let's find out is there any null value present in the dataset.

```
#isNotNull()
```

```
from pyspark.sql.functions import *
#filter data by null values
df.filter(df.name.isNotNull()).show()
```

These PySpark functions are the combination of both the languages Python and SQL.